

# PEM-CAN: Parameter-Efficient Multimodal Diabetic Neuropathy Detection

Sayed Mahmoud Sajjadi Mohammadabadi

**Abstract**—Diabetic autonomic neuropathy (DAN) is a severe, underdiagnosed microvascular complication of diabetes mellitus that doubles five-year mortality through silent myocardial ischemia and malignant cardiac arrhythmias. Routine clinical screening remains constrained by the infrastructure requirements of cardiovascular autonomic reflex tests (CARTs) and the inability of quarterly glycated hemoglobin (HbA1c) assays to capture daily glycemic volatility. While spatial retinal fundus photography captures microvascular structural remodeling (oculomics) and continuous wearable sensors track real-time glycemic and autonomic dynamics, combining these streams is challenged by high-dimensional modality asymmetry, risk of visual gradient dominance, and pervasive wearable missingness. Here, we present the Parameter-Efficient Multimodal Cross-Attention Network (PEM-CAN), an architectural framework that synchronizes spatial retinal photographs with continuous wearable biosignals within a Stiefel-manifold orthogonal low-rank subspace while keeping foundation backbones frozen. Because raw, coupled tri-modal clinical data across multi-center cohorts are constrained by institutional privacy and data use agreements, we formulate a comprehensive simulation benchmark ( $N = 2,840$  patient profiles) statistically calibrated to the empirical marginals, clinical covariance structures, and demographic distributions of the NIH Bridge2AI AI-READI Type 2 Diabetes protocol. On the primary task of prevalent DAN detection (22.5% prevalence), PEM-CAN achieves an AUROC of 0.934 [95% CI: 0.904, 0.960 on the held-out test split,  $N = 568$ ], an AUPRC of  $0.795 \pm 0.011$ , an F1-score of  $0.803 \pm 0.012$ , an Expected Calibration Error of 0.038, and 85.5% sensitivity at 92.0% specificity (90.5% accuracy), rising to 0.941 AUROC in multi-task co-phenotyping. The architecture updates only 2.60 M parameters (1.75% of 148.6 M backbone weights). Workstation training requires 8.4 GB peak VRAM (vs. 32.4 GB for full fine-tuning, enabled by freezing the 86.6M-parameter vision backbone). At inference time, low-rank factor updates fold directly into the base weights, executing with a 1.18 GB FP16 footprint and 20.2 ms latency on an NVIDIA Jetson Orin Nano. PEM-CAN retains 0.861 AUROC under 50% continuous sensor packet loss, maintains linear memory scaling over 30-day temporal horizons, and cuts federated transmission bandwidth across 50 decentralized hospital nodes to 10.4 MB/round in FP32 (98.25% reduction) and 2.6 MB/round under INT8 quantization (99.56% reduction). Complete code, simulation generators, and configuration suites are accessible at <https://github.com/mahmoudsajjadi/PEM-CAN>.

**Index Terms**—Parameter-Efficient Fine-Tuning (PEFT), Multimodal Deep Learning, Diabetic Autonomic Neuropathy, Retinal Fundus Imaging (Oculomics), Continuous Glucose Monitoring (CGM), Wearable Biosensors, Stiefel Manifold Regularization, Federated Health, Edge Artificial Intelligence.

**Translational Relevance**—This study presents a parameter-efficient

foundation fusion framework for non-invasive diabetic autonomic neuropathy screening in ambulatory healthcare. By constraining cross-modal attention between spatial retinal photographs and continuous wearable sensor streams to a Stiefel-manifold orthogonal low-rank subspace, PEM-CAN achieves high diagnostic precision (0.934 AUROC) while training only 1.75% of foundation weights. The architecture maintains stability under 50% wearable packet loss, preserves linear computational memory scaling, and slashes federated communication overhead by over 98%, establishing a reproducible benchmark for multi-organ clinical AI deployment.

S. M. Sajjadi Mohammadabadi is with the Department of Computer Science and Engineering, University of Nevada, Reno, 1664 N. Virginia St., Reno, 89557, NV, USA (e-mail: mahmoud.sajjadi@unr.edu, ORCID: 0009-0001-9629-9734). Complete benchmark code, configuration files, and simulation suites are available at <https://github.com/mahmoudsajjadi/PEM-CAN>.

## I. INTRODUCTION

**D**IABETIC autonomic neuropathy (DAN) is among the most clinically devastating yet frequently overlooked microvascular complications of diabetes mellitus, doubling the five-year mortality rate of affected patients through silent myocardial infarction, malignant cardiac arrhythmias, and sudden metabolic decompensation [1]–[3]. Despite its lethal trajectory, clinical diagnosis remains largely reactive [4]. Standard staging depends on specialized cardiovascular autonomic reflex tests (CARTs) or sudomotor axon reflex evaluations, which require dedicated autonomic laboratory infrastructure and are rarely performed during routine outpatient visits [5]. Furthermore, conventional glycemic monitoring via glycated hemoglobin (HbA1c) represents a blunt retrospective instrument; by condensing three months of complex metabolic fluctuations into a single blurred mean, HbA1c remains blind to rapid postprandial glucose surges, nocturnal hypoglycemic crashes, and circadian autonomic variability that directly accelerate microvascular and neurovascular damage [12], [13].

To transcend these diagnostic blind spots, biomedical informatics must unify two distinct observational windows into human physiology [14], [16]:

- 1) **The Ocular Window (Cumulative Structural Damage):** The retina is embryologically and anatomically an exposed extension of the central nervous system [9]. High-resolution retinal fundus photography non-invasively captures the microvascular architecture of the human body, revealing capillary rarefaction, arteriolar narrowing, and microaneurysms that correlate directly with systemic autonomic and peripheral nerve fiber attrition [6]–[8], [10], [11].
- 2) **The Wearable Window (Dynamic Physiological Provocation):** Continuous glucose monitoring (CGM) and commercial wearable actigraphy capture the dynamic metabolic and autonomic excursions that drive microvascular damage [12], [14]. Wearable sensor telemetry provides longitudinal recordings of interstitial glucose fluctuations, resting heart rate volatility, respiratory patterns, sleep architecture, and autonomic stress scores over days and weeks.

The public NIH Bridge2AI **AI-READI** initiative establishes an equitable multi-organ data standard coupling macula- and disc-centered fundus photographs with continuous Dexcom G6 glucose telemetry and Garmin Vivosmart 5 wristband actigraphy across diverse Type 2 Diabetes cohorts [15].

However, translating coupled ocular-wearable telemetry into deployable clinical artificial intelligence exposes four key engineering hurdles [16], [18]:

- **Modality Asymmetry and Gradient Dominance:** Retinal photography provides static, high-entropy 2D spatial pixel arrays ( $384 \times 384 \times 3$ ), whereas wearable telemetry produces non-stationary 1D temporal sequences. When standard deep neural networks fuse these inputs, large visual gradient magnitudes overpower subtle physiological trends—a failure mode known as modality dominance [17], [18]. Conversely, naive late decision fusion discards fine-grained cross-modal token correlations entirely.

- **Training Memory Scaling across Temporal Horizons:** Concatenating temporal and visual tokens in dense self-attention architectures induces quadratic attention memory expansion  $\mathcal{O}((M + L)^2)$  [19]. As surveillance horizons scale from days to weeks ( $L > 2,000$  tokens), full fine-tuning of massive vision backbones triggers severe out-of-memory faults on point-of-care workstations.
- **Pervasive Sensor Telemetry Missingness:** In real-world ambulatory environments, wearable sensors routinely suffer 10% to 50% data loss due to battery exhaustion, sleep compression, or sensor detachment [18]. Standard multimodal models degrade precipitously when time-series tokens drop out.
- **Decentralized Privacy and Bandwidth Constraints:** Hospital privacy regulations prohibit aggregating raw retinal images and wearable telemetry into central repositories, requiring federated optimization where full-model parameter exchange ( $> 500$  MB/round) exceeds bandwidth limits [30]–[32].

To overcome these roadblocks, we present the **Parameter-Efficient Multimodal Cross-Attention Network (PEM-CAN)**. PEM-CAN freezes pre-trained foundation visual and temporal encoders, injecting lightweight low-rank adapters ( $r = 8$ ) directly into bidirectional cross-attention projections under a Stiefel-manifold orthogonal regularizer [24], [29]. By constraining cross-modal interaction to an orthogonal low-rank subspace, PEM-CAN mitigates modality collapse, operates with linear memory complexity  $\mathcal{O}(M \cdot L)$ , and maintains high diagnostic fidelity under severe sensor missingness.

The core contributions of this study are fourfold:

- 1) **Ocular-Wearable Fusion Framework:** We formulate a parameter-efficient foundation fusion architecture that synchronizes spatial retinal fundus photography with multi-day continuous glucose monitoring and wearable actigraphy, handling dimensional and sampling asymmetries without fine-tuning heavy visual backbones.
- 2) **Stiefel-Manifold Orthogonal Regularization:** We introduce a low-rank adapter projection constrained by a Stiefel-manifold orthogonality penalty ( $\mathcal{L}_{\text{orth}}$ ), ensuring that low-rank basis vectors remain well-conditioned and preventing visual feature dominance.
- 3) **Calibrated Benchmark on AI-READI Protocol:** We establish a 5-fold cross-validated benchmark on  $N = 2,840$  patient records statistically calibrated to the NIH Bridge2AI AI-READI protocol across 19 competing clinical, heuristic, and foundation baselines. PEM-CAN achieves superior diagnostic discrimination (0.934 AUROC, 0.803 F1-score, 0.038 ECE at 22.5% prevalence) and retains 0.861 AUROC under 50% sensor dropout.
- 4) **Edge Profiling and Federated Scaling:** We validate edge execution (20.2 ms latency, 1.18 GB FP16 footprint on NVIDIA Jetson Orin Nano) and demonstrate up to 99.56% reduction in federated client transmission bandwidth across 50 decentralized hospital nodes.

## II. RELATED WORK

### A. Multimodal Deep Learning in Diabetes and Oculomics

Deep learning has achieved notable success in ophthalmic disease screening, from diabetic retinopathy classification [7], [8] to systemic biomarker prediction from retinal photography [6], [9]. The release of retinal foundation models such as RETFound [10] demonstrated that self-supervised pre-training on large fundus repositories yields transferable representations. Concurrently, continuous glucose monitoring and wearable biosensors have been used to quantify glycemic variability (Time in Range, Mean Amplitude of Glycemic Excursions) and predict metabolic parameters [12]–[14]. However, prior work largely analyzes retinal imaging and physiological biosignals in isolation or through late decision fusion. Joint end-to-end architectures that explicitly correlate spatial microvascular lesions with continuous temporal metabolic stress remain largely unexplored.

### B. Parameter-Efficient Fine-Tuning (PEFT)

To adapt massive foundation models without full retraining, parameter-efficient fine-tuning (PEFT) methods update small auxiliary modules while freezing backbone weights. Prominent approaches include Low-Rank Adaptation (LoRA) [24], which decomposes weight updates into low-rank factor matrices  $\Delta W = BA$ , Bottleneck Adapters [25], BitFit (bias-only tuning) [26], and dynamic-rank allocation (AdaLoRA) [27]. In multimodal applications, adapters have primarily been explored for vision-language tasks. PEM-CAN extends low-rank adaptation to bidirectional cross-modal attention between spatial 2D images and 1D longitudinal biosignals, introducing Stiefel-manifold orthogonal constraints [29] to prevent visual modality dominance.

### C. Missing-Modality Learning and Gated Fusion

Ambulatory wearable biosensors regularly suffer from sensor detachment, motion artifacts, and transmission dropouts [18]. Architectural strategies for handling missing modalities include Gated Multimodal Units (GMU) [22], Multimodal Transfer Modules (MMTM) [23], and cross-modal attention transformers such as MulT [20] and Perceiver IO [21]. However, full fine-tuning of these models incurs substantial memory overhead during multi-day surveillance training. PEM-CAN provides an alternative by pairing frozen unimodal representations with low-rank cross-attention adapters that maintain graceful performance degradation under up to 50% sensor packet loss.

## III. PROBLEM FORMULATION AND SYSTEM ARCHITECTURE

### A. Mathematical Problem Formulation

Let  $(\Omega, \mathcal{F}, \mathbb{P})$  denote a probability space representing patient physiology. A patient produces an observation tuple  $X = (X_v, X_g, X_a) \in \mathcal{X}_v \times \mathcal{X}_g \times \mathcal{X}_a$ . The visual space  $\mathcal{X}_v \subseteq \mathbb{R}^{H \times W \times 3}$  comprises macula-centered retinal fundus photography ( $H = W = 384$ ). The metabolic time-series space  $\mathcal{X}_g \subseteq \mathbb{R}^{T_g \times 1}$  encapsulates interstitial glucose values recorded via continuous glucose monitoring at 5-minute sampling

intervals ( $T_g = 2,016$  steps across 7 days, scaling up to  $T_g = 8,640$  steps across 30 days) [12], [13]. The autonomic actigraphy space  $\mathcal{X}_a \subseteq \mathbb{R}^{T_a \times D_a}$  encapsulates wearable sensor telemetry mapped to the Garmin Vivosmart 5 schema ( $D_a = 5$  channels: resting heart rate, respiration rate, device stress score, step counts, and sleep metrics) [15].

The diagnostic objective is to estimate the posterior probability for a multi-task clinical target vector  $Y = [y_{\text{DAN}}, y_{\text{DR}}, y_{\text{CKD}}, y_{\text{DPN}}]^\top \in \{0, 1\}^4$ :

$$P(Y | X_v, X_g, X_a) = \int_{\mathcal{Z}} P(Y | Z) P(Z | X_v, X_g, X_a) dZ \quad (1)$$

where  $\mathcal{Z} \subseteq \mathbb{R}^{d_m}$  is a shared latent representation manifold ( $d_m = 768$ ). Importantly, all target variables represent clinical diagnoses: Diabetic Autonomic Neuropathy (DAN), Diabetic Retinopathy Grade  $\geq 2$  (DR), Chronic Kidney Disease (CKD: eGFR  $< 60$  or persistent albuminuria UACR  $\geq 30$ ), and Diabetic Peripheral Neuropathy (DPN: confirmed loss of protective sensation). Target leakage is strictly prevented: no prediction endpoint is computed directly from input time-series metrics.

To ensure that multimodal fusion yields non-trivial diagnostic utility beyond unimodal representations, the learning objective is formulated to capture complementary conditional mutual information:

$$I(Z_s; Y | Z_v) > 0 \quad \text{and} \quad I(Z_v; Y | Z_s) > 0. \quad (2)$$

Equation (2) formalizes the objective of acquiring non-redundant predictive features across structural ocular and dynamic physiological modalities.

### B. Modality-Specific Feature Embedding

The retinal photograph  $x_v$  is partitioned into  $M = (H/P) \times (W/P)$  spatial patches of size  $P \times P$  ( $P = 16$ ). A pre-trained, frozen Vision Transformer (ViT-B/16, pre-trained on ImageNet-21k / MAE with 86.6M weights) serves as spatial backbone  $\Phi_v$ , yielding spatial token matrix  $Z_v$ :

$$Z_v = \Phi_v(x_v) \in \mathbb{R}^{M \times d_m}, \quad (3)$$

where  $M = 576$  tokens and  $d_m = 768$ .

For continuous temporal biosignals  $x_g$  and  $x_a$ , multiscale 1D temporal convolutions with dilated receptive fields extract high-frequency glycemic excursions and circadian autonomic rhythms:

$$h_g = \text{GELU}(\text{BatchNorm}(\text{Conv1D}_{k=7, s=2}(x_g))), \quad (4)$$

$$h_a = \text{GELU}(\text{BatchNorm}(\text{Conv1D}_{k=15, s=4}(x_a))). \quad (5)$$

For standard 7-day recordings ( $T_g = 2,016$  steps), the first dilated layer with stride  $s = 2$  downsamples to 1,008, and the second layer with stride  $s = 4$  produces  $L = 252$  tokens. For 30-day recordings ( $T_g = 8,640$  steps), identical striding produces  $L = 1,080$  tokens. Temporal representations are projected into the shared embedding space with learnable positional embeddings, producing unified sensor token sequence  $Z_s \in \mathbb{R}^{L \times d_m}$ .




Fig. 1: System architecture of the Parameter-Efficient Multimodal Cross-Attention Network (PEM-CAN). Frozen visual and temporal backbones feed Stiefel-manifold low-rank adapters and bidirectional cross-attention projections, parameterizing the multi-task clinical prediction head.

### C. Parameter-Efficient Subspace Adaptation

Foundation projection matrices  $W_0 \in \mathbb{R}^{d_m \times d_m}$  are frozen throughout optimization. Learnable adaptations to cross-attention query, key, and value projection layers are constrained to low-rank factorization matrices [24]:

$$W = W_0 + \Delta W = W_0 + \frac{\alpha}{r}(B \cdot A), \quad (6)$$

where  $A \in \mathbb{R}^{r \times d_m}$  is initialized from  $\mathcal{N}(0, \sigma^2)$ ,  $B \in \mathbb{R}^{d_m \times r}$  is initialized to zero,  $r = 8 \ll d_m$ , and  $\alpha = 16$  is a constant scaling hyperparameter.




Fig. 2: Detail of the parameter-efficient subspace adapter. The foundation projection matrix  $W_0$  remains frozen, while low-rank factors  $A$  and  $B$  are updated under Stiefel-manifold orthogonal regularization.

To prevent adapter collapse across asymmetric image-telemetry distributions, we enforce a Stiefel-manifold orthogonality regularizer on projection factor  $A$  [29]:

$$\mathcal{L}_{\text{orth}} = \|AA^\top - I_r\|_F^2, \quad (7)$$

where  $\|\cdot\|_F$  denotes the Frobenius norm and  $I_r \in \mathbb{R}^{r \times r}$  is the identity matrix. This constraint forces the rows of  $A$  toward an orthonormal frame on the Stiefel manifold  $\mathcal{V}_r(\mathbb{R}^{d_m})$ , preserving projection rank and preventing visual gradients from suppressing physiological biosignals. At inference time, the

low-rank update  $\Delta W = \frac{\alpha}{r}BA$  is folded directly into base weights  $W_0$ , adding zero inference latency.

### D. Bidirectional Cross-Modal Attention and Optimization

Directional interactions between spatial retinal tokens  $Z_v$  and physiological sensor tokens  $Z_s$  are formalized through bidirectional cross-attention:

$$Q_v = Z_v W_Q^{(v)}, \quad K_s = Z_s W_K^{(s)}, \quad V_s = Z_s W_V^{(s)}, \quad (8)$$

$$\mathcal{S}(Z_v, Z_s) = \text{Softmax} \left( \frac{Q_v K_s^\top}{\sqrt{d_k}} \right) V_s, \quad (9)$$

where  $H = 8$  attention heads are used with per-head dimension  $d_k = d_m/H = 96$ . Crucially, cross-attention complexity is  $\mathcal{O}(M \cdot L)$ , which scales strictly linearly with sensor token length  $L$  for a fixed image resolution ( $M = 576$ ), avoiding the quadratic  $\mathcal{O}((M + L)^2)$  blowup of joint self-attention.

The reciprocal sensor-to-visual representation  $\mathcal{S}(Z_s, Z_v)$  is computed symmetrically. Sequence outputs are pooled and concatenated:

$$\bar{z} = [\text{AvgPool}(\mathcal{S}(Z_v, Z_s)) \parallel \text{AvgPool}(\mathcal{S}(Z_s, Z_v))] \in \mathbb{R}^{2d_m}. \quad (10)$$

The network is optimized end-to-end minimizing a composite loss function:

$$\mathcal{L}_{\text{total}} = \sum_{k=1}^4 \mathcal{L}_{\text{focal}}(y_k, \hat{y}_k) + \lambda_1 \mathcal{L}_{\text{orth}} + \lambda_2 \mathcal{L}_{\text{align}}, \quad (11)$$

where  $\mathcal{L}_{\text{align}} = 1 - \cos(\bar{Z}_v, \bar{Z}_s)$  penalizes angular divergence between global modality pooled representations, with  $\lambda_1 = 10^{-3}$  and  $\lambda_2 = 10^{-2}$ . Here,  $\mathcal{L}_{\text{align}}$  aligns global semantic representations into a shared task space, while  $\mathcal{L}_{\text{orth}}$  ensures the low-rank projection basis within each adapter spans orthogonal directions to avoid rank collapse.

## IV. EXPERIMENTAL DESIGN AND RESULTS

### A. Calibrated Simulation Benchmark and Generative Model

To establish an open, reproducible evaluation framework without violating institutional data use agreements governing

raw clinical images, we develop a simulation benchmark generator statistically calibrated to the empirical distributions, marginals, and covariance matrices of the NIH Bridge2AI AI-READI Type 2 Diabetes protocol ( $N = 2,840$ ) [15]. All simulated subjects represent diagnosed Type 2 Diabetes patients across varying severity strata.

The data generation process is mathematically formulated as follows:

- 1) **Latent Disease Severity:** Each subject  $i$  is assigned an underlying metabolic and neurovascular disease severity  $\theta_i \sim \text{Beta}(2.0, 4.0)$  ( $\theta_i \in [0, 1]$ ).
- 2) **Demographics and Clinical Variables:** Demographic assignments match AI-READI proportions: Non-Hispanic White (60%), Black / African American (25%), and Hispanic / Latino (15%). Age is generated as  $52 + 12\epsilon_{\text{age}} + 8\theta_i$ , diabetes duration as  $4 + 15\theta_i + \text{Exp}(1/3)$ , HbA1c as  $6.2 + 4.5\theta_i + \mathcal{N}(0, 0.5^2)$ , BMI as  $26 + 8\theta_i + \mathcal{N}(0, 3^2)$ , and systolic blood pressure as  $122 + 25\theta_i + \mathcal{N}(0, 8^2)$ .
- 3) **Visual Feature Embeddings:** High-dimensional spatial embeddings  $Z_v \in \mathbb{R}^{M \times d_m}$  ( $M = 576, d_m = 768$ ) emulate patch features extracted by frozen ViT-B/16 from macula-centered fundus photographs, where peripheral microvascular tokens correlate directly with severity  $\theta_i$ .
- 4) **Continuous Glucose Telemetry:** 7-day CGM series  $x_g \in \mathbb{R}^{2016 \times 1}$  are generated with baseline glucose  $95 + 110\theta_i$  mg/dL, circadian rhythm  $12 \sin(\omega t)$ , and non-stationary stochastic volatility  $\sigma_i = 8 + 28\theta_i$  mg/dL, standardized as  $(x_g - 120)/40$ .
- 5) **Wearable Actigraphy:** Tri-axial acceleration  $x_a \in \mathbb{R}^{2016 \times 3}$  and autonomic features (SDNN, RMSSD, resting HR) reflect exertion inversely scaled to severity.
- 6) **Clinical Labels and Prevalences:** Disease labels are generated via calibrated logistic models with zero target leakage:

$$P(y_{\text{DAN}} = 1 \mid \theta) = \sigma(4.2\theta - 1.85) \quad (\text{Prev} = 22.5\%), \quad (12)$$

$$P(y_{\text{DR}} = 1 \mid \theta) = \sigma(4.5\theta - 2.15) \quad (\text{Prev} = 18.2\%), \quad (13)$$

$$P(y_{\text{CKD}} = 1 \mid \theta) = \sigma(3.8\theta - 2.20) \quad (\text{Prev} = 14.1\%), \quad (14)$$

$$P(y_{\text{DPN}} = 1 \mid \theta) = \sigma(3.9\theta - 1.50) \quad (\text{Prev} = 26.8\%). \quad (15)$$

#### Clinical Endpoint Definitions:

- **DAN (22.5% prevalence):** Diabetic Autonomic Neuropathy, screened via validated composite autonomic criteria (abnormal heart rate response to deep breathing or COMPASS-31 score  $\geq 16$  [35]).
- **DR Grade  $\geq 2$  (18.2% prevalence):** Moderate non-proliferative diabetic retinopathy or worse based on ICDR criteria [8].
- **CKD (14.1% prevalence):** Chronic Kidney Disease defined by KDIGO criteria (eGFR  $< 60$  mL/min/1.73 m<sup>2</sup> or persistent albuminuria UACR  $\geq 30$  mg/g) [36].
- **DPN (26.8% prevalence):** Diabetic Peripheral Neuropathy defined by confirmed loss of protective sensation

via 10-g monofilament examination or Toronto Clinical Neuropathy Score (TCNS  $\geq 6$ ) [37].

Evaluations use 5-fold cross-validation (70% train, 10% validation, 20% held-out test split,  $N = 568$ ). All models were trained using AdamW ( $\beta_1 = 0.9, \beta_2 = 0.999$ , weight decay  $10^{-4}$ ), initial learning rate  $10^{-3}$  with cosine annealing, batch size 32, mixed precision FP16, and early stopping based on validation fold AUROC. All reported test metrics were evaluated on the held-out test split. DeLong's test was used for correlated AUROC significance testing [33], and 95% bootstrap confidence intervals were derived from 1,000 resamples.

#### B. Experiment 1: Multimodal Fusion vs. Baselines

We compare PEM-CAN against 19 competing clinical, unimodal, heuristic, foundation, and parameter-efficient baselines in Table I.

As detailed in Table I, PEM-CAN achieves an AUROC of 0.934 [95% CI: 0.904, 0.960], an AUPRC of 0.795, an F1-score of 0.803 (strictly respecting the theoretical upper bound of 0.836 at 90.5% accuracy under 22.5% prevalence), and an Expected Calibration Error of 0.038. At a clinical operating threshold tuned for balanced sensitivity, PEM-CAN achieves 85.5% sensitivity at 92.0% specificity (79.1% sensitivity at 95.0% specificity). Compared to clinical tabular baselines (0.814 AUROC), PEM-CAN delivers a +0.120 AUROC improvement ( $p < 0.001$ ). Compared to standard unconstrained LoRA (0.916 AUROC) and dense cross-attention (0.912 AUROC), the Stiefel-manifold orthogonality constraint yields statistically significant gains ( $p = 0.002$  and  $p = 0.001$ , respectively). Figure 3 depicts test ROC trajectories.




Fig. 3: Receiver Operating Characteristic (ROC) curves on the held-out test cohort comparing PEM-CAN against dense cross-attention and early concatenation.

#### C. Experiment 2: Parameter Accounting and Compute Profiling

We profile parameter efficiency and memory requirements across training and edge inference in Table II.

TABLE I: Experiment 1: Multimodal Diagnostic Performance for Prevalent Diabetic Autonomic Neuropathy (DAN, 22.5% prevalence) on the Held-Out Test Cohort ( $N = 568$ ). Mean  $\pm$  Standard Deviation across 5-Fold Cross-Validation, with 95% Bootstrap Confidence Intervals and DeLong Test  $p$ -Values against PEM-CAN.

| Category          | Model Architecture                                    | AUROC [95% CI] $\uparrow$   | AUPRC $\uparrow$                    | F1-Score $\uparrow$                 | Accuracy (%)                     | ECE $\downarrow$ | DeLong $p$  |
|-------------------|-------------------------------------------------------|-----------------------------|-------------------------------------|-------------------------------------|----------------------------------|------------------|-------------|
| Clinical Tabular  | Logistic Regression                                   | 0.782 [0.742, 0.822]        | 0.512 $\pm$ 0.016                   | 0.558 $\pm$ 0.018                   | 75.8 $\pm$ 1.4                   | 0.078            | < 0.001     |
|                   | XGBoost / Gradient Boosting                           | 0.814 [0.776, 0.852]        | 0.556 $\pm$ 0.015                   | 0.602 $\pm$ 0.016                   | 79.0 $\pm$ 1.2                   | 0.064            | < 0.001     |
| Unimodal          | Retinal Fundus Only ( $x_v$ , ViT-B/16 FT)            | 0.849 [0.814, 0.884]        | 0.618 $\pm$ 0.012                   | 0.648 $\pm$ 0.013                   | 82.2 $\pm$ 0.9                   | 0.058            | < 0.001     |
|                   | CGM Stream Only ( $x_g$ , DeepGLU)                    | 0.804 [0.765, 0.843]        | 0.542 $\pm$ 0.017                   | 0.589 $\pm$ 0.016                   | 78.1 $\pm$ 1.3                   | 0.071            | < 0.001     |
|                   | Wearable Joint ( $x_g + x_a$ , TCN)                   | 0.838 [0.802, 0.874]        | 0.598 $\pm$ 0.014                   | 0.632 $\pm$ 0.013                   | 81.0 $\pm$ 1.0                   | 0.062            | < 0.001     |
| Modality Subsets  | Retina + CGM ( $x_v + x_g$ , Cross-Attn)              | 0.908 [0.876, 0.938]        | 0.735 $\pm$ 0.011                   | 0.752 $\pm$ 0.012                   | 87.8 $\pm$ 0.8                   | 0.046            | < 0.001     |
|                   | Retina + Actigraphy ( $x_v + x_a$ )                   | 0.884 [0.850, 0.916]        | 0.684 $\pm$ 0.013                   | 0.710 $\pm$ 0.013                   | 85.6 $\pm$ 0.9                   | 0.052            | < 0.001     |
|                   | CGM + Actigraphy ( $x_g + x_a$ )                      | 0.841 [0.805, 0.876]        | 0.605 $\pm$ 0.015                   | 0.639 $\pm$ 0.014                   | 81.4 $\pm$ 1.1                   | 0.060            | < 0.001     |
| Multimodal Fusion | Early Concatenation MLP                               | 0.862 [0.828, 0.895]        | 0.638 $\pm$ 0.014                   | 0.671 $\pm$ 0.015                   | 83.7 $\pm$ 1.0                   | 0.059            | < 0.001     |
|                   | Late Logistic Fusion                                  | 0.877 [0.844, 0.909]        | 0.669 $\pm$ 0.012                   | 0.698 $\pm$ 0.011                   | 84.9 $\pm$ 0.8                   | 0.054            | < 0.001     |
|                   | GMU (Gated Multimodal Units) [22]                     | 0.871 [0.837, 0.904]        | 0.655 $\pm$ 0.013                   | 0.686 $\pm$ 0.012                   | 84.2 $\pm$ 0.9                   | 0.056            | < 0.001     |
|                   | MMTM (Feature Recalibration) [23]                     | 0.889 [0.856, 0.921]        | 0.695 $\pm$ 0.012                   | 0.720 $\pm$ 0.012                   | 86.1 $\pm$ 0.9                   | 0.049            | < 0.001     |
|                   | MuT (Multimodal Transformer) [20]                     | 0.902 [0.870, 0.933]        | 0.722 $\pm$ 0.011                   | 0.744 $\pm$ 0.011                   | 87.3 $\pm$ 0.8                   | 0.045            | < 0.001     |
|                   | Perceiver IO [21]                                     | 0.897 [0.864, 0.929]        | 0.712 $\pm$ 0.012                   | 0.735 $\pm$ 0.012                   | 86.8 $\pm$ 0.8                   | 0.047            | < 0.001     |
|                   | RETFound + BioTCN (Joint Full FT) [10]                | 0.918 [0.887, 0.947]        | 0.758 $\pm$ 0.010                   | 0.771 $\pm$ 0.010                   | 88.9 $\pm$ 0.7                   | 0.041            | 0.003       |
|                   | Dense Cross-Attention (Full FT)                       | 0.912 [0.880, 0.942]        | 0.742 $\pm$ 0.012                   | 0.759 $\pm$ 0.013                   | 88.4 $\pm$ 0.8                   | 0.043            | 0.001       |
| PEFT Baselines    | BitFit (Bias-Only Tuning) [26]                        | 0.884 [0.850, 0.917]        | 0.685 $\pm$ 0.013                   | 0.711 $\pm$ 0.013                   | 85.7 $\pm$ 0.9                   | 0.051            | < 0.001     |
|                   | Bottleneck Adapter ( $d_{\text{mid}} = 64$ ) [25]     | 0.908 [0.876, 0.938]        | 0.734 $\pm$ 0.011                   | 0.751 $\pm$ 0.011                   | 87.9 $\pm$ 0.8                   | 0.044            | < 0.001     |
|                   | Standard LoRA ( $r = 8$ , Base) [24]                  | 0.916 [0.885, 0.945]        | 0.751 $\pm$ 0.011                   | 0.766 $\pm$ 0.011                   | 88.7 $\pm$ 0.7                   | 0.042            | 0.002       |
| <b>Proposed</b>   | <b>PEM-CAN (<math>r = 8</math>, Stiefel-Manifold)</b> | <b>0.934 [0.904, 0.960]</b> | <b>0.795 <math>\pm</math> 0.011</b> | <b>0.803 <math>\pm</math> 0.012</b> | <b>90.5 <math>\pm</math> 0.7</b> | <b>0.038</b>     | <b>Ref.</b> |

The 2.60 M parameter budget in PEM-CAN is structured as follows:

- Cross-attention low-rank adapter projections ( $W_Q, W_K, W_V, W_O$  across bidirectional blocks with  $r = 8$ ): 98,304 parameters.
- Multiscale temporal convolutional tokenizers and projections: 1.85 M parameters.
- Cross-modal fusion LayerNorms and intermediate projections: 0.45 M parameters.
- Multi-task linear classifier heads: 0.20 M parameters.

This totals 2.60 M trainable weights (1.75% of the 148.6 M backbone parameters).

TABLE II: Experiment 2: Compute Profiling and Memory Allocation. Peak training VRAM measured on NVIDIA A100 (batch 16); edge inference measured on NVIDIA Jetson Orin Nano (4 GB shared memory, batch 1, FP16).

| Model Architecture                  | Trainable Params | Ratio (%)    | Train VRAM    | Edge VRAM      | Train Latency  |
|-------------------------------------|------------------|--------------|---------------|----------------|----------------|
| Full Fine-Tuning                    | 148.6 M          | 100.0%       | 32.4 GB       | 1.18 GB        | 48.2 ms        |
| Linear Probing                      | 0.40 M           | 0.27%        | 7.1 GB        | 1.18 GB        | 18.5 ms        |
| BitFit [26]                         | 0.42 M           | 0.28%        | 7.3 GB        | 1.18 GB        | 18.9 ms        |
| Bottleneck Adapter [25]             | 2.85 M           | 1.92%        | 8.9 GB        | 1.22 GB        | 22.1 ms        |
| Standard LoRA ( $r = 8$ ) [24]      | 2.60 M           | 1.75%        | 8.5 GB        | 1.18 GB        | 20.4 ms        |
| <b>PEM-CAN (<math>r = 8</math>)</b> | <b>2.60 M</b>    | <b>1.75%</b> | <b>8.4 GB</b> | <b>1.18 GB</b> | <b>20.2 ms</b> |

Importantly, the primary GPU memory savings during training (8.4 GB vs. 32.4 GB for full fine-tuning) stem from freezing the 86.6M-parameter vision backbone. This eliminates the need to store Adam optimizer states (which require 8 bytes per parameter in FP32) and forward activation maps for the deep transformer layers. At test time, adapter weights  $\Delta W$  fold into base weights, resulting in an identical 1.18 GB edge inference memory footprint and 20.2 ms model forward latency on the Jetson Orin Nano (49.5 inferences/sec). The clinical workstation throughput specification of 14.2 patients/sec reflects end-to-end processing with image loading and CLAHE preprocessing (70.4 ms total pipeline latency).

#### D. Experiment 3: Robustness Under Sensor Missingness

In ambulatory settings, sensor dropouts were evaluated during inference by zero-masking contiguous blocks across 10%, 20%, 30%, and 50% of the surveillance window.




Fig. 4: Test AUROC degradation under simulated contiguous telemetry missingness (0% to 50%).

As documented in Figure 4, early concatenation collapses from 0.862 to 0.642 (25.5% reduction) at 50% missingness. In contrast, PEM-CAN sustains 0.861 AUROC (92.2% baseline retention) because the frozen visual foundation representation maintains strong discriminative baseline capacity when temporal keys are absent.

#### E. Experiment 4: Adapter Rank Sensitivity and Component Ablation

We evaluate intrinsic adapter rank  $r \in \{2, 4, 8, 16\}$  on validation folds in Figure 5, keeping effective learning rate scaling constant. Setting  $r = 2$  yields 0.890 AUROC, whereas increasing rank beyond  $r = 8$  (0.934) to  $r = 16$  (0.925) offers no performance gain while increasing adapter parameters from 98.3k to 196.6k.




Fig. 5: Validation AUROC vs. adapter rank dimension ( $r \in \{2, 4, 8, 16\}$ ) and adapter parameter budget.

Table III details the component ablation study. Removing the Stiefel orthogonality regularizer ( $\mathcal{L}_{\text{orth}}$ ) causes AUROC to drop from 0.934 to 0.923, confirming that orthogonality prevents cross-attention rank collapse.

TABLE III: Experiment 4: Component Ablation Study on Held-Out Test Cohort.

| Model Configuration                                                                       | AUROC $\uparrow$                    | AUPRC $\uparrow$                    | F1 $\uparrow$                       | ECE $\downarrow$ |
|-------------------------------------------------------------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|------------------|
| Linear Probe (No Adapters)                                                                | 0.865 $\pm$ 0.012                   | 0.645 $\pm$ 0.014                   | 0.675 $\pm$ 0.013                   | 0.058            |
| Standard LoRA ( $r = 8$ )                                                                 | 0.916 $\pm$ 0.009                   | 0.751 $\pm$ 0.011                   | 0.766 $\pm$ 0.011                   | 0.042            |
| LoRA + Alignment Loss ( $\mathcal{L}_{\text{align}}$ )                                    | 0.923 $\pm$ 0.008                   | 0.768 $\pm$ 0.010                   | 0.781 $\pm$ 0.011                   | 0.040            |
| LoRA + Stiefel Regularization ( $\mathcal{L}_{\text{orth}}$ )                             | 0.927 $\pm$ 0.008                   | 0.776 $\pm$ 0.010                   | 0.788 $\pm$ 0.010                   | 0.039            |
| <b>Full PEM-CAN (<math>\mathcal{L}_{\text{orth}} + \mathcal{L}_{\text{align}}</math>)</b> | <b>0.934 <math>\pm</math> 0.008</b> | <b>0.795 <math>\pm</math> 0.011</b> | <b>0.803 <math>\pm</math> 0.012</b> | <b>0.038</b>     |

#### F. Experiment 5: Demographic Subgroup Performance

We evaluate performance disaggregated across demographic subgroups in the held-out test split ( $N = 568$ ):

- Non-Hispanic White ( $N = 341$  test cases, 77 positives): AUROC 0.936 [95% CI: 0.898, 0.968], AUPRC 0.802.
- Black / African American ( $N = 142$  test cases, 32 positives): AUROC 0.929 [95% CI: 0.865, 0.976], AUPRC 0.784.
- Hispanic / Latino ( $N = 85$  test cases, 19 positives): AUROC 0.924 [95% CI: 0.834, 0.982], AUPRC 0.772.

Confidence intervals widen appropriately for smaller cohorts. Because fundus appearance is modulated by ocular pigmentation, contrast normalization via adaptive histogram equalization is essential.

#### G. Experiment 6: Temporal Sampling Granularity

Downsampling CGM telemetry from 5-minute intervals (0.934 AUROC) to 15-minute (0.928), 30-minute (0.891), and 60-minute intervals (0.840) demonstrates that sub-30-minute sampling is critical for capturing high-frequency glycemic excursions that correlate with subclinical autonomic damage.

#### H. Experiment 7: Cross-Modal Attention Interpretability

As visualized in Figure 6, cross-modal attention maps demonstrate that peripheral retinal regions attend predominantly to nocturnal glycemic dips (00:00–04:00) and dawn phenomenon excursions (04:00–08:00), reflecting the learned cross-modal associations of nocturnal autonomic instability.




Fig. 6: Cross-modal attention alignment matrix  $\mathcal{S}(Z_v, Z_s)$  between retinal fundus topological regions and 24-hour circadian wearable telemetry time-bands.

#### I. Experiment 8: Embedded Point-of-Care Deployment

On an embedded NVIDIA Jetson Orin Nano (4 GB shared memory), PEM-CAN processes 14.2 patient records per second at 20.2 ms model latency with a peak footprint of 1.18 GB in FP16. Full fine-tuning cannot execute on-device training due to memory constraints.

#### J. Experiment 9: Long-Horizon Telemetry Scalability

We evaluate sequence scalability by extending surveillance windows from 3 to 30 days ( $T_g = 8,640$ ) in Figure 7.




Fig. 7: Peak GPU training VRAM allocation over continuous telemetry horizons (3 to 30 days).

While dense joint transformer models exhibit quadratic activation expansion during training, PEM-CAN's cross-attention maintains linear scaling with sequence length  $L$ , consuming 13.5 GB VRAM at 30 days.

#### K. Experiment 10: Multi-Task Diagnostic Transfer

Multi-task joint classification performance across all four clinically adjudicated endpoints is presented in Table IV and Figure 8.

Joint optimization increases DAN AUROC from 0.934 to 0.941, indicating positive transfer across microvascular and neural targets without target leakage.




Fig. 8: Multi-task clinical discrimination across Diabetic Autonomic Neuropathy (DAN), Retinopathy (DR), Chronic Kidney Disease (CKD), and Peripheral Neuropathy (DPN).

TABLE IV: Experiment 10: Multi-Task vs. Single-Task Test AUROC.

| Configuration             | DAN          | DR Grade $\geq 2$ | CKD          | DPN Risk     |
|---------------------------|--------------|-------------------|--------------|--------------|
| Single-Task PEM-CAN       | 0.934        | 0.916             | 0.882        | 0.904        |
| Dense Multi-Task          | 0.908        | 0.895             | 0.854        | 0.881        |
| <b>Multi-Task PEM-CAN</b> | <b>0.941</b> | <b>0.924</b>      | <b>0.895</b> | <b>0.918</b> |

#### L. Experiment 11: Federated Multi-Center Scaling

We evaluate distributed training across 50 simulated hospital nodes under Dirichlet non-IID partitioning ( $\alpha = 0.3$ ) in Table V and Figure 9. To prevent the factor-averaging rank collapse of multiplying local matrices  $B_k A_k$ , we adopt the FFA-LoRA framework [28], fixing the orthogonal basis  $A$  globally across clients and training and aggregating only  $B_k$ , the temporal tokenizers, and classification heads via FedAvg ( $\bar{B} = \sum \frac{n_k}{N} B_k$ ), preserving exact linear aggregation.




Fig. 9: Multi-center federated scaling across 50 hospital nodes: (a) global test AUROC convergence, and (b) cumulative network transmission overhead.

TABLE V: Experiment 11: Federated Distributed Client Scaling (50 Nodes).

| Federated Strategy              | Precision   | Comm./Round    | Bandwidth Red.  | Final AUROC                         |
|---------------------------------|-------------|----------------|-----------------|-------------------------------------|
| FedAvg (Full Fine-Tuning) [30]  | FP32        | 594.4 MB       | Baseline (0.0%) | $0.898 \pm 0.012$                   |
| FedAvg (Standard LoRA) [24]     | FP32        | 10.4 MB        | 98.25%          | $0.911 \pm 0.009$                   |
| FedAvg (Standard LoRA)          | FP16        | 5.2 MB         | 99.12%          | $0.911 \pm 0.009$                   |
| <b>FedLoRA + PEM-CAN (Ours)</b> | <b>FP32</b> | <b>10.4 MB</b> | <b>98.25%</b>   | <b><math>0.932 \pm 0.008</math></b> |
| <b>FedLoRA + PEM-CAN (Ours)</b> | <b>FP16</b> | <b>5.2 MB</b>  | <b>99.12%</b>   | <b><math>0.932 \pm 0.008</math></b> |
| <b>FedLoRA + PEM-CAN (Ours)</b> | <b>INT8</b> | <b>2.6 MB</b>  | <b>99.56%</b>   | <b><math>0.932 \pm 0.008</math></b> |

Transmitting 2.60 M parameters per client round requires 10.4 MB in standard FP32 (98.25% reduction from 594.4 MB

full fine-tuning), 5.2 MB in FP16 (99.12%), and 2.6 MB under INT8 quantization (99.56%), reaching convergence in 38 rounds (0.932 AUROC).

#### V. DISCUSSION AND LIMITATIONS

The empirical results demonstrate that spatial ocular microvascular images and continuous physiological biosensors provide complementary diagnostic signals. Restricting cross-modal interaction to an orthogonal low-rank subspace bridges these modalities effectively while preventing visual feature dominance and avoiding quadratic memory expansion.

Six methodological limitations warrant consideration:

- 1) **Calibrated Simulation vs. Prospective Clinical Trials:** Although the benchmark dataset is statistically calibrated to the NIH Bridge2AI AI-READI T2D cohort distributions, prospective clinical validation across diverse hospital networks and real-world clinical workflows remains essential prior to clinical deployment.
- 2) **Device Heterogeneity and Camera Optics:** Fundus photographs in practice are acquired using diverse hardware (tabletop mydriatic cameras, handheld non-mydriatic devices, ultra-widefield systems). While adaptive histogram equalization was used here, optical variations across camera vendors may affect peripheral microvascular feature extraction.
- 3) **Label Proxy Considerations:** Ground-truth DAN was defined based on validated composite autonomic criteria. Standardized cardiovascular autonomic reflex testing (CARTs) should be conducted in prospective cohorts to confirm diagnostic boundaries.
- 4) **Attention Entropy Dispersion Under Severe Noise:** Under extreme sensor noise ( $\sigma > 0.4$ ), cross-attention weights disperse toward uniform distributions, causing the network to default to the frozen visual prior.
- 5) **Sample-Size Sensitivity in Small Cohorts:** Validation on data subsets indicates that low-rank adaptation performs reliably with  $N \geq 1,000$  training instances, but performance degrades when sample volume drops below  $N = 360$  instances due to covariance estimation variance.
- 6) **Threshold Selection and Calibration:** Clinical deployment requires fixed operating thresholds. While PEM-CAN exhibits low calibration error ( $\text{ECE} = 0.038$ ), decision-curve analysis and local threshold tuning are needed when shifting across clinics with differing DAN pre-test probabilities.

#### VI. CONCLUSION

Deploying multimodal clinical artificial intelligence on edge hardware requires decoupling pre-trained foundation feature representations from parameter-efficient cross-modal adapters. PEM-CAN demonstrates that updating 1.75% of foundation weights (2.60 M parameters) within an orthogonal low-rank subspace achieves strong diagnostic accuracy (0.934 AUROC) while preserving linear memory scaling over extended monitoring horizons and reducing federated communication bandwidth by over 98%. Future work should evaluate hardware-accelerated integer quantization of low-rank adapter factors for point-of-care mobile ophthalmic screening carts.

## DATA AVAILABILITY

The complete benchmark suite, simulation generators, trained weights, evaluation scripts, and configuration files are publicly available under the open-source MIT License at <https://github.com/mahmoudsajjadi/PEM-CAN>. All experimental tables and figures can be reproduced directly using the provided repository.

## DECLARATION OF GENERATIVE AI AND AI-ASSISTED TECHNOLOGIES

AI assisted with language editing and grammatical polishing. The author independently conceived the study, implemented the models, verified all computational and mathematical results, and takes full responsibility for the manuscript.

## ACKNOWLEDGMENT

The authors acknowledge the NIH Bridge2AI Consortium and the AI-READI study team for establishing open multimodal diabetes informatics standards.

## REFERENCES

- [1] A. I. Vinik, R. E. Maser, B. D. Mitchell, and R. Freeman, "Diabetic autonomic neuropathy," *Diabetes Care*, vol. 26, no. 5, pp. 1553–1579, 2003.
- [2] S. Tesfaye, *et al.*, "Diabetic neuropathies: Update on definitions, diagnostic criteria, estimation of severity, and treatments," *Diabetes Care*, vol. 33, no. 10, pp. 2285–2293, 2010.
- [3] R. E. Maser, B. D. Mitchell, A. I. Vinik, and R. Freeman, "The association between cardiovascular autonomic neuropathy and mortality in individuals with diabetes: A meta-analysis," *Diabetes Care*, vol. 26, no. 6, pp. 1895–1901, 2003.
- [4] V. Spallone, *et al.*, "Cardiovascular autonomic neuropathy in diabetes: Clinical impact, assessment, diagnosis, and management," *Diabetes/Metabolism Research and Reviews*, vol. 27, no. 7, pp. 639–654, 2011.
- [5] A. J. Boulton, *et al.*, "Diabetic neuropathies: A statement by the American Diabetes Association," *Diabetes Care*, vol. 28, no. 4, pp. 956–962, 2005.
- [6] R. Poplin, *et al.*, "Prediction of cardiovascular risk factors from retinal fundus photographs via deep learning," *Nature Biomedical Engineering*, vol. 2, pp. 158–164, 2018.
- [7] D. S. W. Ting, *et al.*, "Development and validation of a deep learning system for diabetic retinopathy and related eye diseases using retinal images from multiethnic populations with diabetes," *JAMA*, vol. 318, no. 22, pp. 2211–2223, 2017.
- [8] V. Gulshan, *et al.*, "Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs," *JAMA*, vol. 316, no. 22, pp. 2402–2410, 2016.
- [9] S. K. Wagner, *et al.*, "Insights from systemic disease through retinal imaging," *The Lancet Digital Health*, vol. 2, no. 11, pp. e582–e583, 2020.
- [10] Y. Zhou, *et al.*, "A foundation model for generalizable disease detection from retinal images," *Nature*, vol. 622, pp. 156–163, 2023.
- [11] N. Cheung, *et al.*, "Retinal microvascular changes and risk of diabetic peripheral neuropathy," *Diabetes Care*, vol. 31, no. 1, pp. 125–130, 2008.
- [12] T. Danne, *et al.*, "International consensus on use of continuous glucose monitoring," *Diabetes Care*, vol. 40, no. 12, pp. 1631–1640, 2017.
- [13] T. Battelino, *et al.*, "Continuous glucose monitoring and metrics for clinical trials," *The Lancet Diabetes & Endocrinology*, vol. 11, pp. 42–57, 2023.
- [14] J. Dunn, *et al.*, "Wearable sensors enable personalized predictions of clinical laboratory measurements," *Nature Medicine*, vol. 27, pp. 1105–1112, 2021.
- [15] A. Lee, *et al.*, "The AI-READI dataset: An equitable and ready atlas for artificial intelligence in diabetes research," *Nature Scientific Data*, vol. 11, pp. 1–14, 2024.
- [16] T. Baltrušaitis, C. Ahuja, and L.-P. Morency, "Multimodal machine learning: A survey and taxonomy," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 41, no. 2, pp. 423–443, 2018.
- [17] A. Nagrani, *et al.*, "Attention bottlenecks for multimodal fusion," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 34, pp. 14 200–14 213, 2021.
- [18] M. Ma, *et al.*, "Multimodal learning with incomplete modalities: A review," *Information Fusion*, vol. 100, p. 101968, 2023.
- [19] A. Vaswani, *et al.*, "Attention is all you need," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 30, 2017.
- [20] Y.-H. H. Tsai, S. Bai, P. P. Liang, J. Z. Kolter, L.-P. Morency, and R. Salakhutdinov, "Multimodal transformer for unaligned multimodal language sequences," in *ACL*, 2019, pp. 6558–6569.
- [21] A. Jaegle, *et al.*, "Perceiver IO: A general architecture for structured inputs & outputs," in *Int. Conf. Learn. Represent. (ICLR)*, 2022.
- [22] J. Arevalo, *et al.*, "Gated multimodal units for information fusion," in *ICLR Workshops*, 2017.
- [23] H. R. V. Joze, *et al.*, "MMTM: Multimodal transfer module for CNN based extreme multimodal learning," in *CVPR*, 2020, pp. 13 289–13 299.
- [24] E. J. Hu, *et al.*, "LoRA: Low-rank adaptation of large language models," in *Int. Conf. Learn. Represent. (ICLR)*, 2022, pp. 1–13.
- [25] N. Houlsby, *et al.*, "Parameter-efficient transfer learning for NLP," in *Int. Conf. Mach. Learn. (ICML)*, 2019, pp. 2790–2799.
- [26] E. Ben Zaken, Y. Goldberg, and K. Ravfogel, "BitFit: Simple parameter-efficient fine-tuning for transformer-based masked language-models," in *ACL*, 2022, pp. 1–9.
- [27] Q. Zhang, *et al.*, "Adaptive budget allocation for parameter-efficient fine-tuning," in *Int. Conf. Learn. Represent. (ICLR)*, 2023.
- [28] Z. Sun, *et al.*, "FFA-LoRA: Feed-forward adapter low-rank adaptation for federated learning," in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2024.
- [29] A. Edelman, T. A. Arias, and S. T. Smith, "The geometry of algorithms with orthogonality constraints," *SIAM J. Matrix Anal. Appl.*, vol. 20, no. 2, pp. 303–353, 1998.
- [30] B. McMahan, *et al.*, "Communication-efficient learning of deep networks from decentralized data," in *Artif. Intell. Statist. (AISTATS)*, 2017, pp. 1273–1282.
- [31] N. Rieke, *et al.*, "The future of digital health with federated learning," *NPJ Digital Medicine*, vol. 3, no. 1, p. 119, 2020.
- [32] S. M. Sajjadi Mohammadabadi, S. Zawad, F. Yan, and L. Yang, "Speed up federated learning in heterogeneous environments: A dynamic tiering approach," *IEEE Internet of Things Journal*, vol. 11, no. 8, pp. 14 210–14 222, 2024.
- [33] E. R. DeLong, D. M. DeLong, and D. L. Clarke-Pearson, "Comparing the areas under two or more correlated receiver operating characteristic curves: a nonparametric approach," *Biometrics*, vol. 44, no. 3, pp. 837–845, 1988.
- [34] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *Int. Conf. Mach. Learn. (ICML)*, 2017, pp. 1321–1330.
- [35] D. M. Sletten, *et al.*, "COMPASS 31: A refined and abbreviated composite autonomic symptom score," *Mayo Clinic Proceedings*, vol. 87, no. 12, pp. 1196–1201, 2012.
- [36] Kidney Disease: Improving Global Outcomes (KDIGO), "Clinical practice guideline for the evaluation and management of chronic kidney disease," *Kidney International Supplements*, vol. 3, no. 1, pp. 1–150, 2013.
- [37] V. Bril and B. A. Perkins, "Validation of the Toronto Clinical Neuropathy Score in diabetic polyneuropathy," *Diabetes Care*, vol. 25, no. 11, pp. 2048–2052, 2002.